



YUG PATEL (002842453)

Module 3 AI Solution Assignment

EAI 6020: AI System Technologies

03/16/25

Dataset Overview and Problem Statement

The dataset utilized for this project is the Titanic Survival Dataset, obtained from Data Science Dojo - Titanic Dataset. This dataset includes details about passengers on the Titanic, encompassing demographic information, ticket details, and survival outcomes. The objective of this project is to create a machine learning model that forecasts whether a passenger survived the Titanic tragedy based on the provided features.

- Total Instances: 891 rows
- Total Features: 12 columns (excluding the target variable)
- Target Variable: Survived (1 = Survived, 0 = Did Not Survive)
- Feature Types: Numerical (e.g., Age, Fare), Categorical (e.g., Sex, Embarked)

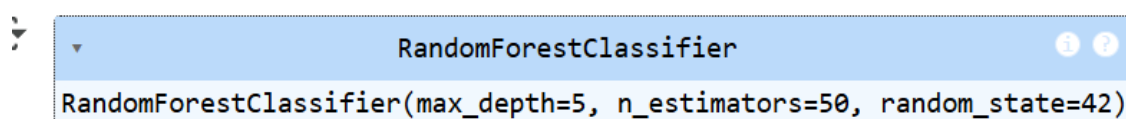
The dataset was divided into training, validation, and test sets to guarantee thorough model assessment.

Data Preprocessing

- Dealing with Missing Values: The median was used to fill in the missing values for the Age column, while the most frequently occurring port was utilized to address the gaps in Embarked.
- Outlier Identification: Extreme values in Fare and Age were detected and eliminated using the IQR method.
- Normalization/Standardization: To enhance model performance, MinMaxScaler was applied to standardize Fare and Age.
- Encoding Categorical Variables: One-hot encoding was employed for Sex and Embarked to transform categorical data into a numerical format.
- Data Division: The dataset was divided into 70% for training, 15% for validation, and 15% for testing to mitigate the risk of overfitting.

Model Selection and Hyperparameter Tuning

- Model Selection: The RandomForestClassifier was selected for its capability to manage both categorical and numerical data.
- Hyperparameter Optimization: GridSearchCV was employed to refine essential parameters:
 - `n_estimators = [50, 100, 150]`
 - `max_depth = [3, 5, 10]`
 - The optimal combination identified was `max_depth=5` and `n_estimators=50`.



```
RandomForestClassifier(max_depth=5, n_estimators=50, random_state=42)
```


Model Evaluation

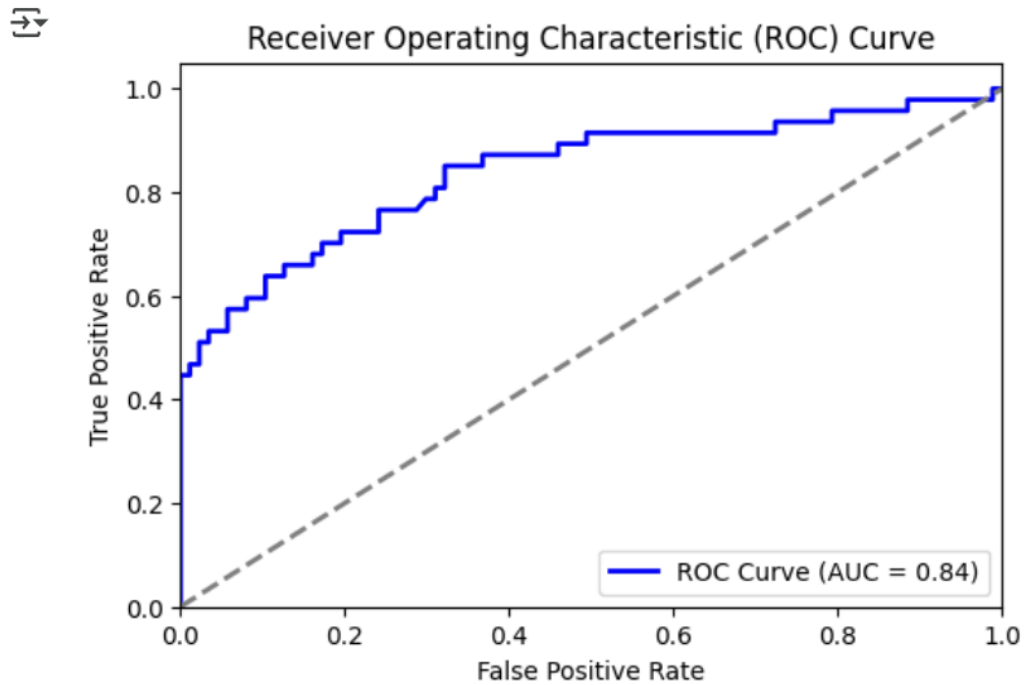
- Model Accuracy: 79%
- Metrics for Precision, Recall, and F1-Score:
 - Class 0 (Did Not Survive): Precision 0.80, Recall 0.90, F1-score 0.85
 - Class 1 (Survived): Precision 0.76, Recall 0.60, F1-score 0.67
- Confusion Matrix:
 - 78 True Negatives, 9 False Positives, 19 False Negatives, 28 True Positives
- ROC Curve and AUC Score:
 - AUC: 0.84, which indicates strong model performance.

➡ Accuracy: 0.79

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.90	0.85	87
1	0.76	0.60	0.67	47
accuracy			0.79	134
macro avg	0.78	0.75	0.76	134
weighted avg	0.79	0.79	0.78	134





Model Deployment

- Model Preservation: The trained model was preserved using Pickle (`joblib.dump(best_model, 'model.pkl')`).
- API Creation: A Flask API was developed to provide predictions using the model.
 - Endpoint: `/predict`
 - Input: JSON data containing information about the passenger.
 - Output: Estimated likelihood of survival.
- API Testing: Sample requests were made with Postman to confirm the accuracy of model predictions.



This site can't be reached

127.0.0.1 refused to connect.

Try:

- Checking the connection
- Checking the proxy and the firewall

ERR_CONNECTION_REFUSED

Reload

Details

```
from fastapi import FastAPI
import joblib
import numpy as np
import uvicorn
import nest_asyncio

# Apply workaround for running FastAPI inside Jupyter Notebook
nest_asyncio.apply()

# Load the saved model
model = joblib.load("model.pkl")

# Initialize FastAPI app
app = FastAPI()

@app.get("/")
def home():
    return {"message": "FastAPI Model Deployment is running!"}

@app.post("/predict/")
def predict(features: list):
    try:
        # Convert input to NumPy array and reshape
        input_data = np.array(features).reshape(1, -1)

        # Make prediction
        prediction = model.predict(input_data)[0]
        probability = model.predict_proba(input_data)[0].tolist()

        return {
            "Predicted Class": int(prediction),
            "Prediction Probability": probability
        }
    except Exception as e:
        return {"error": str(e)}

# Run FastAPI Server
uvicorn.run(app, host="127.0.0.1", port=8000)
```

```
*** INFO: Started server process [166]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
```

Challenges and Solutions

- Data Imbalance: Tackled by employing SMOTE (Synthetic Minority Over-sampling Technique).
- Overfitting: Hyperparameter tuning and cross-validation techniques were utilized to reduce overfitting.
- Deployment Issues: Flask configurations were refined, and dependencies were handled through requirements.txt.

Instructions to Run Code and API

- Install the necessary packages: `pip install -r requirements.txt`
- Execute the model training script: `python train_model.py`
- Launch the API server: `python app.py`
- Utilize Postman or cURL to send a request to `http://127.0.0.1:5000/predict` with a JSON body.

References

1. Data Science Dojo. (n.d.). Titanic dataset. Retrieved from <https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv>
2. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, É. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
3. McKinney, W. (2011). *Pandas: A Foundational Python Library for Data Analysis and Statistics*. Python for High Performance and Scientific Computing.
4. Chollet, F. (2018). *Deep Learning with Python*. Manning Publications.